

A Comparative Analysis and Implementation on Gaussian Naive Bayes (GNB)

Janet Joseph, Jimmy Lin, Paneri Patel, Zehui Zhou

Team - Learning not Machine

Contents

- Introduction
- Methodology
- Implementation
- Results
- Conclusion

What is Naive Bayes?

Naive Bayes is a **probabilistic classifier** based on Bayes' theorem:

$$P(Y | X) = \frac{P(X | Y) P(Y)}{P(X)}$$

Conditional independence assumption: given the class ($Y = c$), all features (x_1, \dots, x_d) are independent.

$$P(X | Y = c) = \prod_{j=1}^d P(x_j | Y = c)$$

Why "Gaussian"?

For continuous features, Gaussian Naive Bayes assumes each feature follows a normal (Gaussian) distribution under each class:

$$x_j \mid Y = c \sim \mathcal{N}(\mu_{jc}, \sigma_{jc}^2)$$

GNB Training Pipeline

1. Estimate priors:

The prior probability of each class (c) is:

$$P(Y = c) = \frac{\text{count of class } c}{N}$$

2. Estimate mean and variance for every feature (j) and every class (c).

$$\text{Mean } (\mu_{c,j}) = \frac{1}{N_c} \sum_{i=1}^{N_c} x_{i,j}$$

$$\text{Variance } (\sigma_{c,j}^2) = \frac{1}{N_c - 1} \sum_{i=1}^{N_c} (x_{i,j} - \mu_{c,j})^2$$

3. Prediction:

Compute the **posterior probability** for each class: $P(Y = c | X) \propto P(Y = c) \prod_{j=1}^d P(x_j | Y = c)$

Then **choose the class with the maximum posterior probability.**

Representation

Gaussian Naive Bayes converts input features into a class prediction by computing the **posterior probability** for each class:

$$P(Y = c \mid X = x_1, \dots, x_d) \propto P(Y = c) \prod_{j=1}^d \mathcal{N}(x_j \mid \mu_{jc}, \sigma_{jc}^2)$$

The classifier predicts the class with the highest posterior. In practice, we work with **log probabilities** (to avoid numerical underflow):

$$\hat{y} = \arg \max_c \left[\log P(Y = c) + \sum_{j=1}^d \log \mathcal{N}(x_j ; \mu_{jc}, \sigma_{jc}^2) \right]$$

We use **log probabilities for numerical stability** because multiplying many small Gaussian likelihoods can lead to extremely tiny numbers that computers cannot represent reliably. Summing logs avoids this problem.

Loss Function

Naive Bayes does not minimize a traditional loss like MSE or cross-entropy using gradient descent.

Instead, the model is trained by **maximizing the likelihood** of the data.

Equivalently, the loss is the **negative log-likelihood**:

$$L = - \sum_{i=1}^N \log P(y^{(i)} | x^{(i)})$$

Gaussian Naive Bayes maximizes the likelihood under the assumption that each feature is normally distributed for each class.

Maximum Likelihood Estimation (MLE). For class c and feature j, the likelihood is modeled as:

$$P(x_j | Y = c) = \frac{1}{\sqrt{2\pi\sigma_{c,j}^2}} \exp\left(-\frac{(x_j - \mu_{c,j})^2}{2\sigma_{c,j}^2}\right)$$

Because the parameters have **closed-form MLE solutions**, no gradient descent, no iterative optimization is required during training.

Optimizer

Gaussian Naive Bayes does **not** use an iterative optimizer such as gradient descent.

Gaussian Naive Bayes has **closed-form Maximum Likelihood Estimates (MLE)** for all parameters.

Prior probabilities:

Gaussian parameters for each feature (j) and class (c):

$$\text{Mean } (\mu_{c,j}) = \frac{1}{N_c} \sum_{i=1}^{N_c} x_{i,j}$$

$$\text{Variance } (\sigma_{c,j}^2) = \frac{1}{N_c - 1} \sum_{i=1}^{N_c} (x_{i,j} - \mu_{c,j})^2$$

These are computed **directly from the data** using simple MLE formulas.

No optimization loop is required.

Code Overview

- Train the model
- Predict outcomes
- Evaluate accuracy
- Model test results

Train the Model

This method computes, for each class:

- Prior probabilities based on class frequencies
- The mean of each feature
- The variance of each feature

X_train : numpy.ndarray

A 2D array of shape (n_examples, n_attributes) containing continuous features.

y_train : numpy.ndarray

A 1D array of shape (n_examples,) containing integer class labels in {0, ..., n_classes-1}.

Pseudo-code

1. Convert X_train and y_train to arrays
2. If X_train is not 2D, raise error
3. If y_train is not 1D, raise error
4. If number of rows in X_train \neq length of y_train, raise error
5. n_examples \leftarrow number of rows in X_train
n_attributes \leftarrow number of columns in X_train
6. Initialize:
 - means \leftarrow empty list
 - vars \leftarrow empty list
 - priors \leftarrow empty list
7. epsilon \leftarrow var_smoothing \times max(variance of X_train across features)
8. For each class c = 0 to n_classes - 1:
 - a. X_c \leftarrow all rows of X_train where y_train = c
 - b. If X_c is empty, raise error
 - c. prior_c \leftarrow |X_c| / n_examples
 - Append prior_c to priors
 - d. mu_c \leftarrow mean of X_c across features
 - e. var_c \leftarrow variance of X_c across features + epsilon
 - f. Append mu_c to means
 - g. Append var_c to vars
9. means \leftarrow array(means)
vars \leftarrow array(vars)
priors \leftarrow array(priors)
10. Return the trained model

Gaussian Probability Function

[_gaussian_pdf](#): Computes Gaussian pdf value for a scalar x.

@params:

x: float

mu: float (mean)

var: float (variance)

@return:

pdf value at x

Pseudo-code

1. $\text{coeff} \leftarrow 1 / \sqrt(2\pi\sigma^2)$
2. $\text{exp_term} \leftarrow \exp(- (x - \mu)^2 / (2\sigma^2))$
3. $\text{pdf} \leftarrow \text{coeff} \times \text{exp_term}$
4. Return pdf

Predict Outcomes

- **predict**: Outputs a predicted label for each input in inputs. Uses log-space to avoid overflow/underflow.

@params:

inputs: a 2D NumPy array containing inputs

@return:

a 1D numpy array of predictions

Pseudo-code

1. Convert inputs to array
2. If inputs is 1D, reshape to 2D
3. If inputs is not 2D, raise error
4. If number of features in inputs \neq number of trained features, raise error
5. Initialize predictions \leftarrow empty list
6. For each input vector x in inputs:
 - a. Initialize log_joint_probs \leftarrow empty list
 - b. For each class $c = 0$ to $n_{\text{classes}} - 1$:
 - i. $\log_prob \leftarrow \log(\text{prior}[c])$
 - ii. For each feature j :
$$\text{pdf} \leftarrow \text{GaussianPDF}(x[j], \text{mean}[c][j], \text{var}[c][j])$$
$$\log_prob \leftarrow \log_prob + \log(\text{pdf} + \epsilon)$$
 - iii. Append \log_prob to \log_joint_probs
 - c. $\text{predicted_label} \leftarrow$ index of maximum value in \log_joint_probs
 - d. Append predicted_label to predictions
7. Return predictions as an array

Predict Outcomes

- **predict_proba**: Outputs posterior probabilities for each class.

@params:

inputs: a 2D NumPy array containing inputs

@return:

a 2D numpy array of probabilities (n_examples x n_classes)

Pseudo-code

1. Convert inputs to array
2. If inputs is 1D, reshape to 2D
3. Initialize probas \leftarrow empty list
4. For each input vector x in inputs:
 - a. Initialize log_joint_probs \leftarrow empty list
 - b. For each class $c = 0$ to $n_classes - 1$:
 - i. $\log_prob \leftarrow \log(\text{prior}[c])$
 - ii. For each feature j :
 $\text{pdf} \leftarrow \text{GaussianPDF}(x[j], \text{mean}[c][j], \text{var}[c][j])$
 $\log_prob \leftarrow \log_prob + \log(\text{pdf} + \epsilon)$
 - iii. Append \log_prob to \log_joint_probs
 - c. Convert log values to probabilities:
 - i. $\max_log \leftarrow \max(\log_joint_probs)$
 - ii. $\exp_shifted \leftarrow \exp(\log_joint_probs - \max_log)$
 - iii. $\text{probs} \leftarrow \exp_shifted / \sum(\exp_shifted)$
 - d. Append probs to probas
5. Return probas as an array

Loss Function

loss: Computes average negative log-likelihood (NLL) loss.

@params:

X: a 2D numpy array of examples

y: a 1D numpy array of labels

@return:

float, average negative log-likelihood

Pseudo-code

1. Convert X and y to arrays
2. proba \leftarrow PredictProba(X)
3. $\epsilon \leftarrow$ small constant
4. Initialize log_probs \leftarrow empty list
5. For each index $i = 0$ to $\text{length}(y) - 1$:
 - a. true_class $\leftarrow y[i]$
 - b. $\log_p \leftarrow \log(\text{proba}[i, \text{true_class}] + \epsilon)$
 - c. Append \log_p to log_probs
6. $\text{loss} \leftarrow -\text{mean}(\log_{\text{probs}})$
7. Return loss

Evaluate Accuracy

accuracy: Outputs the accuracy of the trained model on a given dataset.

score: Returns the mean accuracy on the given test data and labels.

Pseudo-code (accuracy)

@params:

X_test: a 2D numpy array of examples
y_test: a 1D numpy array of labels

@return:

float, accuracy between 0 and 1

1. predictions \leftarrow Predict(X_test)
2. correct \leftarrow number of indices where predictions = y_test
3. accuracy \leftarrow correct / length(y_test)
4. Return accuracy

Pseudo-code (score)

@params:

X: a 2D numpy array of examples
y: a 1D numpy array of true labels

@return:

float, accuracy score

1. score \leftarrow Accuracy(X, y)
2. Return score


Custom GNB vs Sklearn: Statistical Results and Plots






- Prediction Agreement
 - Percentage of identical predictions between models
- Accuracy Results
 - Percentage of correct predictions
- Confidence Statistics
 - Mean confidence: average of max(predict_proba)
- Probability Differences
 - Differences in probability outputs between models
- Timing Statistics
 - How long training takes
- Accuracy Comparison Bar Plot
 - Train vs test accuracy for both implementations
- Confusion Matrix Comparison
 - Where predictions go right/wrong in a grid format
- Training Time Comparison
 - Bar chart of training speed with error bars
- Decision Boundaries (2D only)
 - Colored regions where each class "wins" in feature space

GNB Test Result - *Iris* dataset

We compared our implementation with the scikit-learn Gaussian Naive Bayes model using *Iris* datasets provided by scikit-learn.

```
Statistical summary - Iris:  
Prediction Agreement: 1.0000  
Custom GNB - Train: 0.9810, Test: 0.9111  
Sklearn GNB - Train: 0.9810, Test: 0.9111  
Accuracy Difference: 0.000000  
  
Confidence Statistics:  
Custom GNB Mean Confidence: 0.9698  
Sklearn GNB Mean Confidence: 0.9698  
Custom GNB Correct Prediction Mean Confidence: 0.9800  
Custom GNB Wrong Prediction Mean Confidence: 0.8656  
  
Timing Statistics (10 runs):  
Custom GNB Avg Training Time: 0.000105s ± 0.000316s  
Sklearn GNB Avg Training Time: 0.000632s ± 0.000520s  
Speed Ratio (Custom/Sklearn): 0.17x
```





GNB Test Result - *Breast Cancer* dataset

We compared our implementation with the scikit-learn Gaussian Naive Bayes model using *Breast Cancer* datasets provided by scikit-learn.

Statistical summary - Breast Cancer:

Prediction Agreement: 1.0000

Custom GNB - Train: 0.9422, Test: 0.9474

Sklearn GNB - Train: 0.9447, Test: 0.9474

Accuracy Difference: 0.000000

Confidence Statistics:

Custom GNB Mean Confidence: 0.9871

Sklearn GNB Mean Confidence: 0.9871

Custom GNB Correct Prediction Mean Confidence: 0.9906


Custom GNB Wrong Prediction Mean Confidence: 0.9240



Timing Statistics (10 runs):

Custom GNB Avg Training Time: 0.000300s ± 0.000458s

Sklearn GNB Avg Training Time: 0.000424s ± 0.000523s

Speed Ratio (Custom/Sklearn): 0.71x





Conclusion

- We implemented the Gaussian Naive Bayes model by computing the class-specific means and variances from the training data.
- Using these parameters, we constructed the Gaussian likelihood for each class and used it to compute posterior probabilities for prediction.
- Our implementation achieves the similar accuracy as the scikit-learn Gaussian Naive Bayes classifier across the datasets we tested.
- *Challenges:* Maintaining numerical stability throughout the probability computations and ensuring correct tensor shapes and feature alignment across training and prediction

Q & A

